

# Edge Computing Lab

**Class: TY-AIEC**

**School of Computing, MIT Art Design Technology University**

*Academic Year: 2024-25*

## **Experiment No. 10**

### **Introduction**

Study of Transfer Learning (Images) on Edge Computing Devices

**Objective:** Build a project to apply Transfer Learning of MobileNetV1 & V2 architectures trained on an ImageNet dataset

### **Tasks:**

- Understand Transfer learning
- Understanding of MobileNetV1 & V2 Architectures
- Configure Edge Impulse for Object Detection
- Apply a pre-trained network for you to fine-tune your specific application
- Building and Training a Model
- Deploy on Edge Computing Devices

### **Introduction**

Edge Impulse is a development platform for machine learning on edge devices, targeted at developers who want to create intelligent device solutions. The " Camera "sensor reading equivalent in Edge Impulse would typically involve creating a simple machine learning model that can run on an edge device, like classifying sensor data or recognizing a basic pattern.

### **Materials Required**

- Nano BLE Sense Board

### **Theory**

GPIO (General Purpose Input/Output) pins on the Raspberry Pi are used for interfacing with other electronic components. BCM numbering refers to the pin numbers in the Broadcom SOC channel, which is a more consistent way to refer to the GPIO pins across different versions of the

Here's a high-level overview of steps you'd follow to create a "Hello World" project on Edge Impulse:

### **Steps to Configure the Edge Impulse:**

1. Create an Account and New Project:
  - Sign up for an Edge Impulse account.
  - Create a new project from the dashboard.
2. Connect a Device:

- You can use a supported development board or your smartphone as a sensor device.
- Follow the instructions to connect your device to your Edge Impulse project.

### 3. Collect Data:

- Use the Edge Impulse mobile app or the Web interface to collect data from the onboard sensors.
- For a "Hello World" project, you could collect accelerometer data, for instance.

### 4. Create an Impulse:

- Go to the 'Create impulse' page.
- Add a processing block (e.g., time-series data) and a learning block (e.g., classification).
- Save the impulse, which defines the machine learning pipeline.

### 5. Design a Neural Network:

- Navigate to the 'NN Classifier' under the 'Learning blocks'.
- Design a simple neural network. Edge Impulse provides a default architecture that works well for most basic tasks.

### 6. Train the Model:

- Click on the 'Start training' button to train your machine learning model with the collected data.

### 7. Test the Model:

- Once the model is trained, you can test its performance with new data in the 'Model Testing' tab.

### 8. Deploy the Model:

- Go to the 'Deployment' tab.
- Select the deployment method that suits your edge device (e.g., Arduino library, WebAssembly, container, etc.).
- Follow the instructions to deploy the model to your device.

## 9. Run Inference:

- With the model deployed, run inference on the edge device to see it classifying data in real-time.

## 10. Monitor:

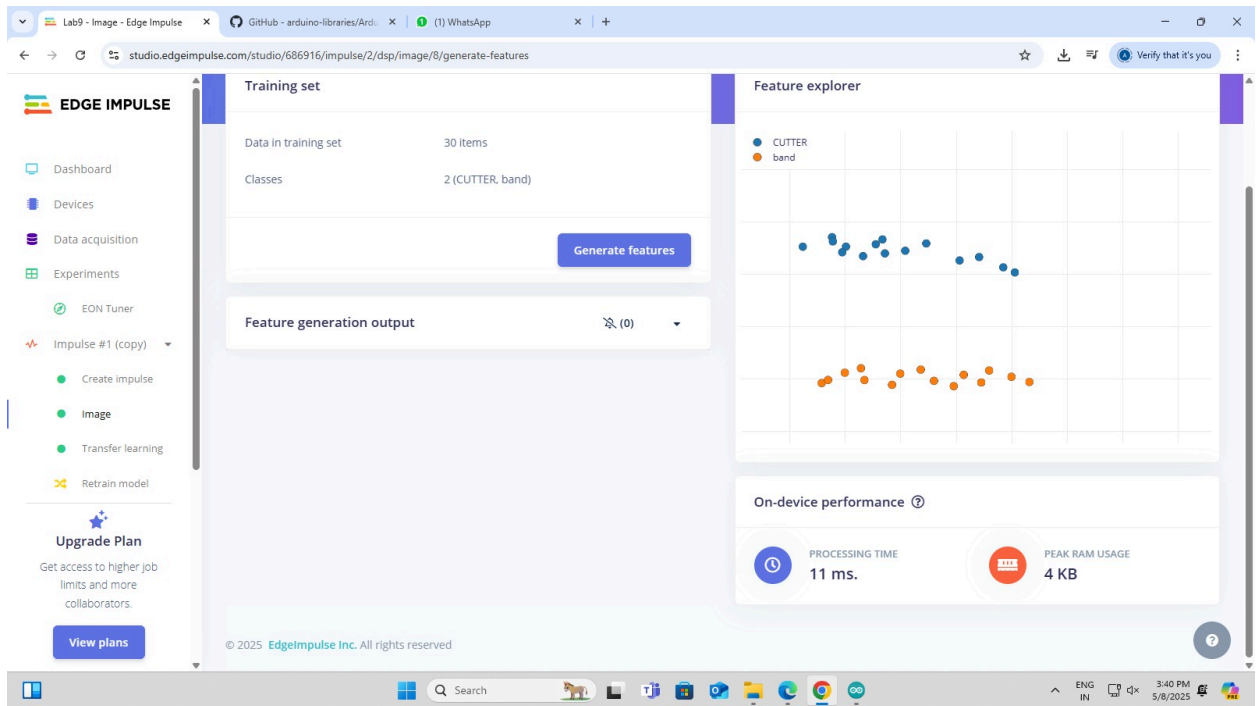
- You can monitor the performance of your device through the Edge Impulse studio.

### 1) Dataset Image

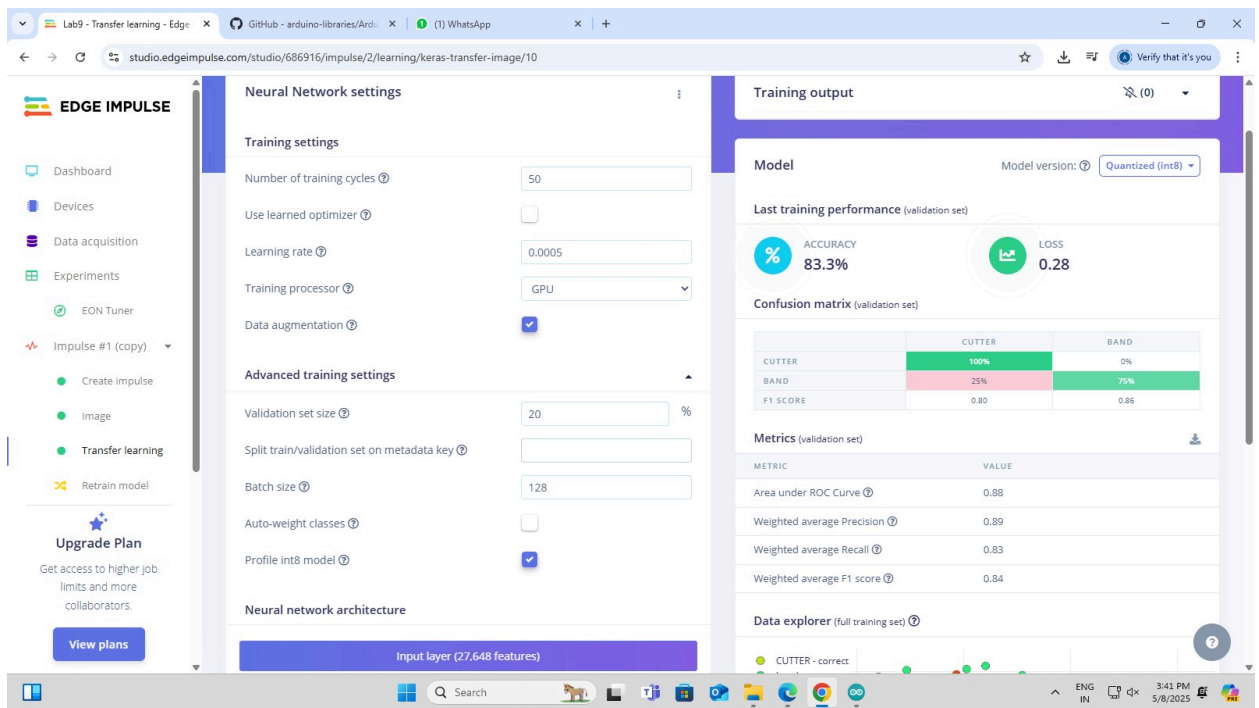
The screenshot displays the Edge Impulse Studio web interface. The top navigation bar includes links for Dataset, Data explorer, Data sources, Synthetic data, AI labeling, and CSV Wizard. The main content area shows a dataset with 40 items collected, a 75% / 25% train/test split, and a 'Collect data' button. Below this, a table lists training samples with columns for Sample Name, Label, and Added date. A 'RAW DATA' section on the right prompts the user to click on a sample to load it.

SAMPLE NAME	LABEL	ADDED
CUTTER.5qk8dopd	CUTTER	Today, 15:01:47
CUTTER.5qk8dfe0	CUTTER	Today, 15:01:38
CUTTER.5qk8d65q	CUTTER	Today, 15:01:28
CUTTER.5qk8csq7	CUTTER	Today, 15:01:19
CUTTER.5qk8cjei	CUTTER	Today, 15:01:09
CUTTER.5qk8ca51	CUTTER	Today, 15:01:00
CUTTER.5qk8c0q0	CUTTER	Today, 15:00:50

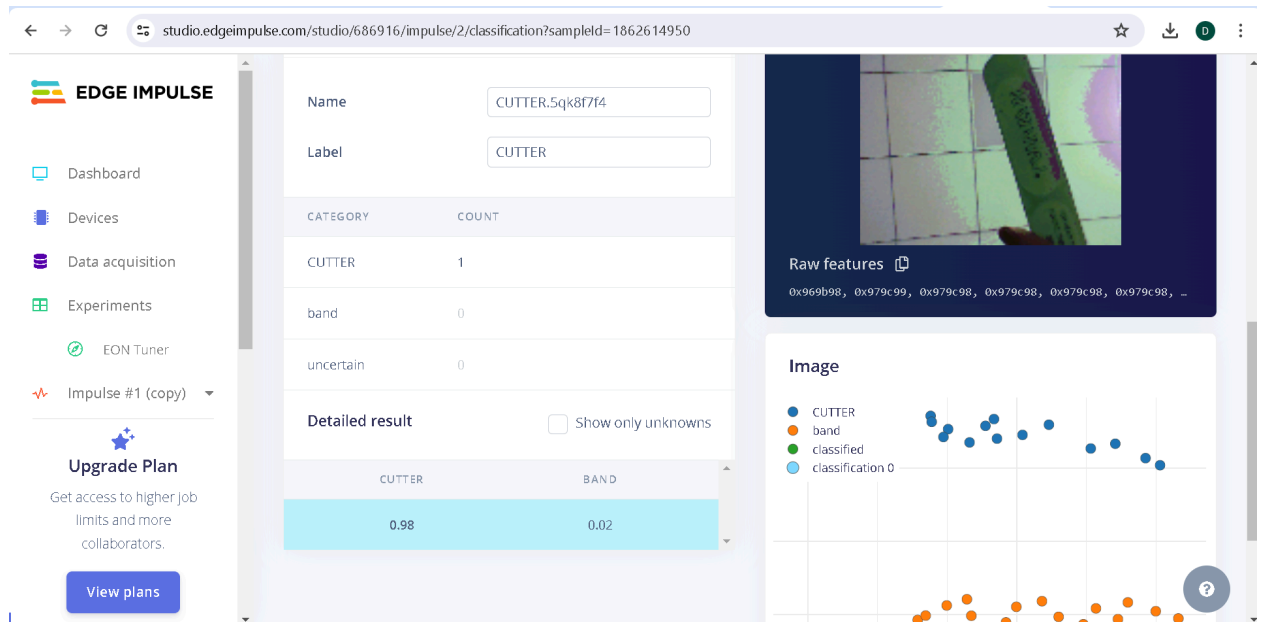
### 2) Feature extraction - Image



### 3) Accuracy / Loss - Confusion Matrix – image



### 4) Validation Result – Image



## 5) Copy the code of Arduino Sketch

```
/* Edge Impulse ingestion SDK

 * Copyright (c) 2022 EdgeImpulse Inc.
 *
 * Licensed under the Apache License, Version 2.0 (the "License");
 * you may not use this file except in compliance with the License.
 * You may obtain a copy of the License at
 *
 * http://www.apache.org/licenses/LICENSE-2.0
 *
 * Unless required by applicable law or agreed to in writing,
software
 * distributed under the License is distributed on an "AS IS" BASIS,
 * WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or
implied.
 * See the License for the specific language governing permissions
and
 * limitations under the License.
 *
 */
```

```

/* Includes
----- */

#include <Lab9_inferencing.h>

#include <Arduino_OV767X.h> //Click here to get the library:
https://www.arduino.cc/reference/en/libraries/arduino\_ov767x/

#include <stdint.h>

#include <stdlib.h>

/* Constant variables
----- */

#define EI_CAMERA_RAW_FRAME_BUFFER_COLS      160

#define EI_CAMERA_RAW_FRAME_BUFFER_ROWS      120

#define DWORD_ALIGN_PTR(a)    ((a & 0x3) ?(((uintptr_t)a + 0x4) &
~(uintptr_t)0x3) : a)

/*
** NOTE: If you run into TFLite arena allocation issue.
**
** This may be due to may dynamic memory fragmentation.
** Try defining "-DEI_CLASSIFIER_ALLOCATION_STATIC" in
boards.local.txt (create
** if it doesn't exist) and copy this file to
**
`<ARDUINO_CORE_INSTALL_PATH>/arduino/hardware/<mbed_core>/<core_vers
ion>/`.
**
** See
**
(https://support.arduino.cc/hc/en-us/articles/360012076960-Where-are-the-installed-cores-located-)

```

```

** to find where Arduino installs cores on your machine.

**

** If the problem persists then there's not enough memory for this
model and application.

*/

/* Edge Impulse
----- */

class OV7675 : public OV767X {

public:

    int begin(int resolution, int format, int fps);

    void readFrame(void* buffer);

private:

    int vsyncPin;

    int hrefPin;

    int pclkPin;

    int xclkPin;

    volatile uint32_t* vsyncPort;

    uint32_t vsyncMask;

    volatile uint32_t* hrefPort;

    uint32_t hrefMask;

    volatile uint32_t* pclkPort;

    uint32_t pclkMask;

    uint16_t width;

    uint16_t height;

    uint8_t bytes_per_pixel;

```

```

        uint16_t bytes_per_row;

        uint8_t buf_rows;

        uint16_t buf_size;

        uint8_t resize_height;

        uint8_t *raw_buf;

        void *buf_mem;

        uint8_t *interp_buf;

        uint8_t *buf_limit;

        void readBuf();

        int allocate_scratch_buffs();

        int deallocate_scratch_buffs();
};

typedef struct {

    size_t width;

    size_t height;
} ei_device_resize_resolutions_t;

/**
 * @brief      Check if new serial data is available
 *
 * @return     Returns number of available bytes
 */
int ei_get_serial_available(void) {

    return Serial.available();
}

```



```

/**
 * @brief      Get next available byte
 *
 * @return     byte
 */
char ei_get_serial_byte(void) {
    return Serial.read();
}

/* Private variables
----- */

static OV7675 Cam;

static bool is_initialised = false;

/*
** @brief points to the output of the capture
*/
static uint8_t *ei_camera_capture_out = NULL;

uint32_t resize_col_sz;

uint32_t resize_row_sz;

bool do_resize = false;

bool do_crop = false;

static bool debug_nn = false; // Set this to true to see e.g.
features generated from the raw signal

/* Function definitions
----- */

bool ei_camera_init(void);

```

```

void ei_camera_deinit(void);

bool ei_camera_capture(uint32_t img_width, uint32_t img_height,
uint8_t *out_buf) ;

int calculate_resize_dimensions(uint32_t out_width, uint32_t
out_height, uint32_t *resize_col_sz, uint32_t *resize_row_sz, bool
*do_resize);

void resizeImage(int srcWidth, int srcHeight, uint8_t *srcImage, int
dstWidth, int dstHeight, uint8_t *dstImage, int iBpp);

void cropImage(int srcWidth, int srcHeight, uint8_t *srcImage, int
startX, int startY, int dstWidth, int dstHeight, uint8_t *dstImage,
int iBpp);

/**
 * @brief      Arduino setup function
 */
void setup()
{
    // put your setup code here, to run once:

    Serial.begin(115200);

    // comment out the below line to cancel the wait for USB
connection (needed for native USB)

    while (!Serial);

    Serial.println("Edge Impulse Inferencing Demo");

    // summary of inferencing settings (from model_metadata.h)

    ei_printf("Inferencing settings:\n");

    ei_printf("\tImage resolution: %dx%d\n",
EI_CLASSIFIER_INPUT_WIDTH, EI_CLASSIFIER_INPUT_HEIGHT);

    ei_printf("\tFrame size: %d\n",
EI_CLASSIFIER_DSP_INPUT_FRAME_SIZE);

```

```

        ei_printf("\tNo. of classes: %d\n",
sizeof(ei_classifier_inferencing_categories) /
sizeof(ei_classifier_inferencing_categories[0]));
    }

/**
 * @brief      Get data and run inferencing
 *
 * @param[in]  debug  Get debug info if true
 */
void loop()
{
    bool stop_inferencing = false;

    while(stop_inferencing == false) {

        ei_printf("\nStarting inferencing in 2 seconds...\n");

        // instead of wait_ms, we'll wait on the signal, this allows
        threads to cancel us...

        if (ei_sleep(2000) != EI_IMPULSE_OK) {
            break;
        }

        ei_printf("Taking photo...\n");

        if (ei_camera_init() == false) {
            ei_printf("ERR: Failed to initialize image sensor\r\n");
            break;
        }
    }
}

```

```

        // choose resize dimensions

        uint32_t resize_col_sz;

        uint32_t resize_row_sz;

        bool do_resize = false;

        int res =
calculate_resize_dimensions(EI_CLASSIFIER_INPUT_WIDTH,
EI_CLASSIFIER_INPUT_HEIGHT, &resize_col_sz, &resize_row_sz,
&do_resize);

        if (res) {

            ei_printf("ERR: Failed to calculate resize dimensions
(%d)\r\n", res);

            break;

        }

        void *snapshot_mem = NULL;

        uint8_t *snapshot_buf = NULL;

        snapshot_mem = ei_malloc(resize_col_sz*resize_row_sz*2);

        if(snapshot_mem == NULL) {

            ei_printf("failed to create snapshot_mem\r\n");

            break;

        }

        snapshot_buf = (uint8_t
*)DWORD_ALIGN_PTR((uintptr_t) snapshot_mem);

        if (ei_camera_capture(EI_CLASSIFIER_INPUT_WIDTH,
EI_CLASSIFIER_INPUT_HEIGHT, snapshot_buf) == false) {

            ei_printf("Failed to capture image\r\n");

            if (snapshot_mem) ei_free(snapshot_mem);

            break;

        }

```

```

    ei::signal_t signal;

    signal.total_length = EI_CLASSIFIER_INPUT_WIDTH *
EI_CLASSIFIER_INPUT_HEIGHT;

    signal.get_data = &ei_camera_cutout_get_data;

    // run the impulse: DSP, neural network and the Anomaly
algorithm

    ei_impulse_result_t result = { 0 };

    EI_IMPULSE_ERROR ei_error = run_classifier(&signal, &result,
debug_nn);

    if (ei_error != EI_IMPULSE_OK) {

        ei_printf("Failed to run impulse (%d)\n", ei_error);

        ei_free(snapshot_mem);

        break;

    }

    // print the predictions

    ei_printf("Predictions (DSP: %d ms., Classification: %d ms.,
Anomaly: %d ms.): \n",

               result.timing.dsp, result.timing.classification,
result.timing.anomaly);

#ifdef EI_CLASSIFIER_OBJECT_DETECTION == 1

    ei_printf("Object detection bounding boxes:\r\n");

    for (uint32_t i = 0; i < result.bounding_boxes_count; i++) {

        ei_impulse_result_bounding_box_t bb =
result.bounding_boxes[i];

        if (bb.value == 0) {

            continue;

        }

    }

```

```

        ei_printf("  %s (%f) [ x: %u, y: %u, width: %u, height:
%u ]\r\n",

                bb.label,

                bb.value,

                bb.x,

                bb.y,

                bb.width,

                bb.height);

    }

    // Print the prediction results (classification)
#else

    ei_printf("Predictions:\r\n");

    for (uint16_t i = 0; i < EI_CLASSIFIER_LABEL_COUNT; i++) {

        ei_printf("  %s: ",
ei_classifier_inferencing_categories[i]);

        ei_printf("%.5f\r\n", result.classification[i].value);

    }

#endif

    // Print anomaly result (if it exists)
#if EI_CLASSIFIER_HAS_ANOMALY

    ei_printf("Anomaly prediction: %.3f\r\n", result.anomaly);

#endif

#if EI_CLASSIFIER_HAS_VISUAL_ANOMALY

    ei_printf("Visual anomalies:\r\n");

    for (uint32_t i = 0; i < result.visual_ad_count; i++) {

```

```

        ei_impulse_result_bounding_box_t bb =
result.visual_ad_grid_cells[i];

        if (bb.value == 0) {

            continue;

        }

        ei_printf("  %s (%f) [ x: %u, y: %u, width: %u, height:
%u ]\r\n",

                    bb.label,

                    bb.value,

                    bb.x,

                    bb.y,

                    bb.width,

                    bb.height);

    }

#endif

    while (ei_get_serial_available() > 0) {

        if (ei_get_serial_byte() == 'b') {

            ei_printf("Inferencing stopped by user\r\n");

            stop_inferencing = true;

        }

    }

    if (snapshot_mem) ei_free(snapshot_mem);

}

ei_camera_deinit();

}

/**

 * @brief      Determine whether to resize and to which dimension

```

```

*

* @param[in]  out_width      width of output image
* @param[in]  out_height    height of output image

* @param[out] resize_col_sz  pointer to frame buffer's
column/width value

* @param[out] resize_row_sz  pointer to frame buffer's
rows/height value

* @param[out] do_resize      returns whether to resize (or not)

*

*/

int calculate_resize_dimensions(uint32_t out_width, uint32_t
out_height, uint32_t *resize_col_sz, uint32_t *resize_row_sz, bool
*do_resize)
{
    size_t list_size = 2;

    const ei_device_resize_resolutions_t list[list_size] = {
{42,32}, {128,96} };

    // (default) conditions

    *resize_col_sz = EI_CAMERA_RAW_FRAME_BUFFER_COLS;

    *resize_row_sz = EI_CAMERA_RAW_FRAME_BUFFER_ROWS;

    *do_resize = false;

    for (size_t ix = 0; ix < list_size; ix++) {

        if ((out_width <= list[ix].width) && (out_height <=
list[ix].height)) {

            *resize_col_sz = list[ix].width;

            *resize_row_sz = list[ix].height;

            *do_resize = true;

            break;

        }
    }
}

```



```

    }

    return 0;
}

/**
 * @brief    Setup image sensor & start streaming
 *
 * @retval   false if initialisation failed
 */
bool ei_camera_init(void) {
    if (is_initialised) return true;

    if (!Cam.begin(QQVGA, RGB565, 1)) { // VGA downsampled to QQVGA
(OV7675)

        ei_printf("ERR: Failed to initialize camera\r\n");

        return false;
    }

    is_initialised = true;

    return true;
}

/**
 * @brief    Stop streaming of sensor data
 */
void ei_camera_deinit(void) {
    if (is_initialised) {
        Cam.end();
    }
}

```

```

        is_initialised = false;

    }

}

/**
 * @brief      Capture, rescale and crop image
 *
 * @param[in]  img_width      width of output image
 * @param[in]  img_height     height of output image
 * @param[in]  out_buf        pointer to store output image, NULL may
be used
 *
 *                                when full resolution is expected.
 *
 * @retval      false if not initialised, image captured, rescaled or
cropped failed
 *
 */
bool ei_camera_capture(uint32_t img_width, uint32_t img_height,
uint8_t *out_buf)
{
    if (!is_initialised) {
        ei_printf("ERR: Camera is not initialized\r\n");
        return false;
    }

    if (!out_buf) {
        ei_printf("ERR: invalid parameters\r\n");
        return false;
    }
}

```

```

    // choose resize dimensions

    int res = calculate_resize_dimensions(img_width, img_height,
&resize_col_sz, &resize_row_sz, &do_resize);

    if (res) {

        ei_printf("ERR: Failed to calculate resize dimensions
(%d)\r\n", res);

        return false;

    }

    if ((img_width != resize_col_sz)
        || (img_height != resize_row_sz)) {

        do_crop = true;

    }

    Cam.readFrame(out_buf); // captures image and resizes

    if (do_crop) {

        uint32_t crop_col_sz;

        uint32_t crop_row_sz;

        uint32_t crop_col_start;

        uint32_t crop_row_start;

        crop_row_start = (resize_row_sz - img_height) / 2;

        crop_col_start = (resize_col_sz - img_width) / 2;

        crop_col_sz = img_width;

        crop_row_sz = img_height;

        //ei_printf("crop cols: %d, rows: %d\r\n",
crop_col_sz, crop_row_sz);

```

```

        cropImage(resize_col_sz, resize_row_sz,

                  out_buf,

                  crop_col_start, crop_row_start,

                  crop_col_sz, crop_row_sz,

                  out_buf,

                  16);

    }

    // The following variables should always be assigned
    // if this routine is to return true
    // cutout values
    //ei_camera_snapshot_is_resized = do_resize;
    //ei_camera_snapshot_is_cropped = do_crop;
    ei_camera_capture_out = out_buf;

    return true;
}

/**
 * @brief      Convert RGB565 raw camera buffer to RGB888
 *
 * @param[in]   offset      pixel offset of raw buffer
 * @param[in]   length      number of pixels to convert
 * @param[out]  out_buf     pointer to store output image
 */
int ei_camera_cutout_get_data(size_t offset, size_t length, float
*out_ptr) {

    size_t pixel_ix = offset * 2;

    size_t bytes_left = length;

```

```

size_t out_ptr_ix = 0;

// read byte for byte
while (bytes_left != 0) {

    // grab the value and convert to r/g/b

    uint16_t pixel = (ei_camera_capture_out[pixel_ix] << 8) |
ei_camera_capture_out[pixel_ix+1];

    uint8_t r, g, b;

    r = ((pixel >> 11) & 0x1f) << 3;

    g = ((pixel >> 5) & 0x3f) << 2;

    b = (pixel & 0x1f) << 3;

    // then convert to out_ptr format

    float pixel_f = (r << 16) + (g << 8) + b;

    out_ptr[out_ptr_ix] = pixel_f;

    // and go to the next pixel

    out_ptr_ix++;

    pixel_ix+=2;

    bytes_left--;

}

// and done!

return 0;
}

// This include file works in the Arduino environment
// to define the Cortex-M intrinsics
#ifdef __ARM_FEATURE_SIMD32

```

```

#include <device.h>

#endif

// This needs to be < 16 or it won't fit. Cortex-M4 only has SIMD
for signed multiplies

#define FRAC_BITS 14

#define FRAC_VAL (1<<FRAC_BITS)

#define FRAC_MASK (FRAC_VAL - 1)

//

// Resize

//

// Assumes that the destination buffer is dword-aligned

// Can be used to resize the image smaller or larger

// If resizing much smaller than 1/3 size, then a more robust
algorithm should average all of the pixels

// This algorithm uses bilinear interpolation - averages a 2x2
region to generate each new pixel

//

// Optimized for 32-bit MCUs

// supports 8 and 16-bit pixels

void resizeImage(int srcWidth, int srcHeight, uint8_t *srcImage, int
dstWidth, int dstHeight, uint8_t *dstImage, int iBpp)
{
    uint32_t src_x_accum, src_y_accum; // accumulators and fractions
for scaling the image

    uint32_t x_frac, nx_frac, y_frac, ny_frac;

    int x, y, ty, tx;

    if (iBpp != 8 && iBpp != 16)

        return;

    src_y_accum = FRAC_VAL/2; // start at 1/2 pixel in to account
for integer downsampling which might miss pixels

```

```

const uint32_t src_x_frac = (srcWidth * FRAC_VAL) / dstWidth;

const uint32_t src_y_frac = (srcHeight * FRAC_VAL) / dstHeight;

const uint32_t r_mask = 0xf800f800;

const uint32_t g_mask = 0x07e007e0;

const uint32_t b_mask = 0x001f001f;

uint8_t *s, *d;

uint16_t *s16, *d16;

uint32_t x_frac2, y_frac2; // for 16-bit SIMD

for (y=0; y < dstHeight; y++) {

    ty = src_y_accum >> FRAC_BITS; // src y

    y_frac = src_y_accum & FRAC_MASK;

    src_y_accum += src_y_frac;

    ny_frac = FRAC_VAL - y_frac; // y fraction and 1.0 - y
fraction

    y_frac2 = ny_frac | (y_frac << 16); // for M4/M4 SIMD

    s = &srcImage[ty * srcWidth];

    s16 = (uint16_t *)&srcImage[ty * srcWidth * 2];

    d = &dstImage[y * dstWidth];

    d16 = (uint16_t *)&dstImage[y * dstWidth * 2];

    src_x_accum = FRAC_VAL/2; // start at 1/2 pixel in to
account for integer downsampling which might miss pixels

    if (iBpp == 8) {

        for (x=0; x < dstWidth; x++) {

            uint32_t tx, p00,p01,p10,p11;

            tx = src_x_accum >> FRAC_BITS;

            x_frac = src_x_accum & FRAC_MASK;

            nx_frac = FRAC_VAL - x_frac; // x fraction and 1.0 - x
fraction

            x_frac2 = nx_frac | (x_frac << 16);

```

```

        src_x_accum += src_x_frac;

        p00 = s[tx]; p10 = s[tx+1];

        p01 = s[tx+srcWidth]; p11 = s[tx+srcWidth+1];

#ifdef __ARM_FEATURE_SIMD32

        p00 = __SMLAD(p00 | (p10<<16), x_frac2, FRAC_VAL/2) >>
FRAC_BITS; // top line

        p01 = __SMLAD(p01 | (p11<<16), x_frac2, FRAC_VAL/2) >>
FRAC_BITS; // bottom line

        p00 = __SMLAD(p00 | (p01<<16), y_frac2, FRAC_VAL/2) >>
FRAC_BITS; // combine

#else // generic C code

        p00 = ((p00 * nx_frac) + (p10 * x_frac) + FRAC_VAL/2) >>
FRAC_BITS; // top line

        p01 = ((p01 * nx_frac) + (p11 * x_frac) + FRAC_VAL/2) >>
FRAC_BITS; // bottom line

        p00 = ((p00 * ny_frac) + (p01 * y_frac) + FRAC_VAL/2) >>
FRAC_BITS; // combine top + bottom

#endif // Cortex-M4/M7

        *d++ = (uint8_t)p00; // store new pixel

    } // for x

} // 8-bpp

else

{ // RGB565

    for (x=0; x < dstWidth; x++) {

        uint32_t tx, p00,p01,p10,p11;

        uint32_t r00, r01, r10, r11, g00, g01, g10, g11, b00,
b01, b10, b11;

        tx = src_x_accum >> FRAC_BITS;

        x_frac = src_x_accum & FRAC_MASK;

        nx_frac = FRAC_VAL - x_frac; // x fraction and 1.0 - x
fraction

```



```

        x_frac2 = nx_frac | (x_frac << 16);

        src_x_accum += src_x_frac;

        p00 = __builtin_bswap16(s16[tx]); p10 =
__builtin_bswap16(s16[tx+1]);

        p01 = __builtin_bswap16(s16[tx+srcWidth]); p11 =
__builtin_bswap16(s16[tx+srcWidth+1]);

#ifdef __ARM_FEATURE_SIMD32
    {

        p00 |= (p10 << 16);

        p01 |= (p11 << 16);

        r00 = (p00 & r_mask) >> 1; g00 = p00 & g_mask; b00 = p00
& b_mask;

        r01 = (p01 & r_mask) >> 1; g01 = p01 & g_mask; b01 = p01
& b_mask;

        r00 = __SMLAD(r00, x_frac2, FRAC_VAL/2) >> FRAC_BITS; //
top line

        r01 = __SMLAD(r01, x_frac2, FRAC_VAL/2) >> FRAC_BITS; //
bottom line

        r00 = __SMLAD(r00 | (r01<<16), y_frac2, FRAC_VAL/2) >>
FRAC_BITS; // combine

        g00 = __SMLAD(g00, x_frac2, FRAC_VAL/2) >> FRAC_BITS; //
top line

        g01 = __SMLAD(g01, x_frac2, FRAC_VAL/2) >> FRAC_BITS; //
bottom line

        g00 = __SMLAD(g00 | (g01<<16), y_frac2, FRAC_VAL/2) >>
FRAC_BITS; // combine

        b00 = __SMLAD(b00, x_frac2, FRAC_VAL/2) >> FRAC_BITS; //
top line

        b01 = __SMLAD(b01, x_frac2, FRAC_VAL/2) >> FRAC_BITS; //
bottom line

        b00 = __SMLAD(b00 | (b01<<16), y_frac2, FRAC_VAL/2) >>
FRAC_BITS; // combine

    }

```

```

    #else // generic C code

        {

            r00 = (p00 & r_mask) >> 1; g00 = p00 & g_mask; b00 = p00
& b_mask;

            r10 = (p10 & r_mask) >> 1; g10 = p10 & g_mask; b10 = p10
& b_mask;

            r01 = (p01 & r_mask) >> 1; g01 = p01 & g_mask; b01 = p01
& b_mask;

            r11 = (p11 & r_mask) >> 1; g11 = p11 & g_mask; b11 = p11
& b_mask;

            r00 = ((r00 * nx_frac) + (r10 * x_frac) + FRAC_VAL/2) >>
FRAC_BITS; // top line

            r01 = ((r01 * nx_frac) + (r11 * x_frac) + FRAC_VAL/2) >>
FRAC_BITS; // bottom line

            r00 = ((r00 * ny_frac) + (r01 * y_frac) + FRAC_VAL/2) >>
FRAC_BITS; // combine top + bottom

            g00 = ((g00 * nx_frac) + (g10 * x_frac) + FRAC_VAL/2) >>
FRAC_BITS; // top line

            g01 = ((g01 * nx_frac) + (g11 * x_frac) + FRAC_VAL/2) >>
FRAC_BITS; // bottom line

            g00 = ((g00 * ny_frac) + (g01 * y_frac) + FRAC_VAL/2) >>
FRAC_BITS; // combine top + bottom

            b00 = ((b00 * nx_frac) + (b10 * x_frac) + FRAC_VAL/2) >>
FRAC_BITS; // top line

            b01 = ((b01 * nx_frac) + (b11 * x_frac) + FRAC_VAL/2) >>
FRAC_BITS; // bottom line

            b00 = ((b00 * ny_frac) + (b01 * y_frac) + FRAC_VAL/2) >>
FRAC_BITS; // combine top + bottom

        }

    #endif // Cortex-M4/M7

    r00 = (r00 << 1) & r_mask;

    g00 = g00 & g_mask;

    b00 = b00 & b_mask;

```

```

        p00 = (r00 | g00 | b00); // re-combine color components

        *d16++ = (uint16_t) __builtin_bswap16(p00); // store new
pixel

        } // for x

        } // 16-bpp

    } // for y
} /* resizeImage() */

//

// Crop

//

// Assumes that the destination buffer is dword-aligned

// optimized for 32-bit MCUs

// Supports 8 and 16-bit pixels

//

void cropImage(int srcWidth, int srcHeight, uint8_t *srcImage, int
startX, int startY, int dstWidth, int dstHeight, uint8_t *dstImage,
int iBpp)
{

    uint32_t *s32, *d32;

    int x, y;

    if (startX < 0 || startX >= srcWidth || startY < 0 || startY >=
srcHeight || (startX + dstWidth) > srcWidth || (startY + dstHeight)
> srcHeight)

        return; // invalid parameters

    if (iBpp != 8 && iBpp != 16)

        return;

    if (iBpp == 8) {

        uint8_t *s, *d;

```

```

for (y=0; y<dstHeight; y++) {

    s = &srcImage[srcWidth * (y + startY) + startX];

    d = &dstImage[(dstWidth * y)];

    x = 0;

    if ((intptr_t)s & 3 || (intptr_t)d & 3) { // either src or
dst pointer is not aligned

        for (; x<dstWidth; x++) {

            *d++ = *s++; // have to do it byte-by-byte

        }

    } else {

        // move 4 bytes at a time if aligned or alignment not
enforced

        s32 = (uint32_t *)s;

        d32 = (uint32_t *)d;

        for (; x<dstWidth-3; x+= 4) {

            *d32++ = *s32++;

        }

        // any remaining stragglers?

        s = (uint8_t *)s32;

        d = (uint8_t *)d32;

        for (; x<dstWidth; x++) {

            *d++ = *s++;

        }

    }

} // for y

} // 8-bpp

else

{

    uint16_t *s, *d;

```

```

    for (y=0; y<dstHeight; y++) {

        s = (uint16_t *)&srcImage[2 * srcWidth * (y + startY) +
startX * 2];

        d = (uint16_t *)&dstImage[(dstWidth * y * 2)];

        x = 0;

        if ((intptr_t)s & 2 || (intptr_t)d & 2) { // either src or
dst pointer is not aligned

            for (; x<dstWidth; x++) {

                *d++ = *s++; // have to do it 16-bits at a time

            }

        } else {

            // move 4 bytes at a time if aligned or alignment no
enforced

            s32 = (uint32_t *)s;

            d32 = (uint32_t *)d;

            for (; x<dstWidth-1; x+= 2) { // we can move 2 pixels at a
time

                *d32++ = *s32++;

            }

            // any remaining stragglers?

            s = (uint16_t *)s32;

            d = (uint16_t *)d32;

            for (; x<dstWidth; x++) {

                *d++ = *s++;

            }

        }

    } // for y

} // 16-bpp case

} /* cropImage() */

```

```

#if !defined(EI_CLASSIFIER_SENSOR) || EI_CLASSIFIER_SENSOR !=
EI_CLASSIFIER_SENSOR_CAMERA

#error "Invalid model for current sensor"

#endif

// OV767X camera library override

#include <Arduino.h>

#include <Wire.h>

#define digitalPinToBitMask(P) (1 << (digitalPinToPinName(P) % 32))

#define portInputRegister(P) ((P == 0) ? &NRF_P0->IN : &NRF_P1->IN)

//

// OV7675::begin()

//

// Extends the OV767X library function. Some private variables are
needed

// to use the OV7675::readFrame function.

//

int OV7675::begin(int resolution, int format, int fps)
{

    pinMode(OV7670_VSYNC, INPUT);

    pinMode(OV7670_HREF, INPUT);

    pinMode(OV7670_PLK, INPUT);

    pinMode(OV7670_XCLK, OUTPUT);

    vsyncPort = portInputRegister(digitalPinToPort(OV7670_VSYNC));

    vsyncMask = digitalPinToBitMask(OV7670_VSYNC);

    hrefPort = portInputRegister(digitalPinToPort(OV7670_HREF));

```

```

hrefMask = digitalPinToBitMask(OV7670_HREF);

pclkPort = portInputRegister(digitalPinToPort(OV7670_PLK));

pclkMask = digitalPinToBitMask(OV7670_PLK);


// init driver to use full image sensor size
bool ret = OV767X::begin(VGA, format, fps);

width = OV767X::width(); // full sensor width
height = OV767X::height(); // full sensor height
bytes_per_pixel = OV767X::bytesPerPixel();

bytes_per_row = width * bytes_per_pixel; // each pixel is 2
bytes

resize_height = 2;


buf_mem = NULL;

raw_buf = NULL;

intrp_buf = NULL;

//allocate_scratch_buffs();


return ret;
} /* OV7675::begin() */


int OV7675::allocate_scratch_buffs()
{

    //ei_printf("allocating buffers..\r\n");

    buf_rows = height / resize_row_sz * resize_height;

    buf_size = bytes_per_row * buf_rows;


    buf_mem = ei_malloc(buf_size);

    if(buf_mem == NULL) {

```

```

        ei_printf("failed to create buf_mem\r\n");

        return false;

    }

    raw_buf = (uint8_t *)DWORD_ALIGN_PTR((uintptr_t)buf_mem);

    //ei_printf("allocating buffers OK\r\n");

    return 0;
}

int OV7675::deallocate_scratch_buffs()
{
    //ei_printf("deallocating buffers...\r\n");

    ei_free(buf_mem);

    buf_mem = NULL;

    //ei_printf("deallocating buffers OK\r\n");

    return 0;
}

//
// OV7675::readFrame()
//
// Overrides the OV767X library function. Fixes the camera output to
be
// a far more desirable image. This image utilizes the full sensor
size
// and has the correct aspect ratio. Since there is limited memory
on the
// Nano we bring in only part of the entire sensor at a time and
then

```



```

// interpolate to a lower resolution.

//

void OV7675::readFrame(void* buffer)
{
    allocate_scratch_buffs();

    uint8_t* out = (uint8_t*)buffer;

    noInterrupts();

    // Falling edge indicates start of frame

    while ((*vsyncPort & vsyncMask) == 0); // wait for HIGH
    while ((*vsyncPort & vsyncMask) != 0); // wait for LOW

    int out_row = 0;

    for (int raw_height = 0; raw_height < height; raw_height +=
buf_rows) {

        // read in 640xbuf_rows buffer to work with

        readBuf();

        resizeImage(width, buf_rows,

                    raw_buf,

                    resize_col_sz, resize_height,

                    &(out[out_row]),

                    16);

        out_row += resize_col_sz * resize_height * bytes_per_pixel;
/* resize_col_sz * 2 * 2 */

    }
}

```

```

interrupts();

deallocate_scratch_buffs();
} /* OV7675::readFrame() */

//
// OV7675::readBuf()
//
// Extends the OV767X library function. Reads buf_rows VGA rows from
the
// image sensor.
//
void OV7675::readBuf()
{
    int offset = 0;

    uint32_t ulPin = 33; // P1.xx set of GPIO is in 'pin' 32 and
above

    NRF_GPIO_Type * port;

    port = nrf_gpio_pin_port_decode(&ulPin);

    for (int i = 0; i < buf_rows; i++) {
        // rising edge indicates start of line

        while ((*hrefPort & hrefMask) == 0); // wait for HIGH

        for (int col = 0; col < bytes_per_row; col++) {
            // rising edges clock each data byte

            while ((*pclkPort & pclkMask) != 0); // wait for LOW

```

```

        uint32_t in = port->IN; // read all bits in parallel

        in >>= 2; // place bits 0 and 1 at the "bottom" of the
register

        in &= 0x3f03; // isolate the 8 bits we care about

        in |= (in >> 6); // combine the upper 6 and lower 2 bits

        raw_buf[offset++] = in;

        while ((*pclkPort & pclkMask) == 0); // wait for HIGH
    }

    while ((*hrefPort & hrefMask) != 0); // wait for LOW
}

} /* OV7675::readBuf() */

```

## 6) Screen shot of Arduino Terminal - Result

